

DeepTCR-based Prediction of CD19 CAR-T Cell Response in Refractory/Relapsed B-cell Lymphoma: A Deep Learning Approach Using Supervised and Unsupervised Analysis

Mateenah Jahan¹, Sherry Bhalla^{1,*†}

¹CSIR-Institute of Genomics and Integrative Biology (CSIR-IGIB), New Delhi, India

*Corresponding author: sherry.igib@csir.res.in

†Research Supervisor

Research Trainee Final Submission – CSIR-IGIB

January 11, 2026

Abstract

Background: Chimeric antigen receptor T-cell (CAR-T) therapy targeting CD19 has revolutionized the treatment of refractory/relapsed B-cell lymphoma, yet response rates remain variable with only 40–60% of patients achieving durable remission. T-cell receptor (TCR) repertoires of infused CAR-T products encode information about T-cell fitness and therapeutic potential, but traditional repertoire metrics have shown limited predictive value. Deep learning approaches offer potential to extract complex patterns from high-dimensional TCR data for clinical prediction.

Methods: We analyzed 239,637 TCR sequences from 34 patients with refractory/relapsed B-cell lymphoma treated with CD19 CAR-T therapy (18 responders, 16 non-responders). Using DeepTCR, we applied both **supervised** and **unsupervised** deep learning approaches. The supervised model employs attention-based multiple instance learning (MIL) with 100-fold Monte Carlo cross-validation to predict treatment response, integrating CDR3 amino acid sequences with V-gene and J-gene usage patterns. The unsupervised model uses variational autoencoder (VAE)-based clustering to discover TCR sequence patterns and repertoire structure without label information. Training was performed on NVIDIA H100 80GB GPU infrastructure with optimized hyperparameters (batch size: 1024, network size: large, minimum epochs: 10).

Results: The **supervised** deep learning model achieved robust predictive performance with a mean area under the receiver operating characteristic curve (AUC) of 0.754 ± 0.035 (SD) across 100 Monte Carlo cross-validation folds (range: 0.658–0.794, median: 0.765, 95% CI: 0.747–0.761 from 1,000 bootstrap iterations). The attention mechanism identified 2,379 high-attention sequences (top 1%), with 72% of the top 100 predictive sequences originating from responder patients. Enrichment analysis revealed preferential usage of TRBV5-8 (1.54-fold) and TRBJ2-2 (1.17-fold) in high-attention sequences. The **unsupervised** VAE analysis revealed distinct TCR clustering patterns that separate responders from non-responders

based on repertoire structure, identifying shared sequence motifs associated with CAR-T product fitness. Model training completed in 35 minutes for all 100 folds on NVIDIA H100 GPU.

Conclusions: TCR repertoire analysis using both supervised and unsupervised deep learning provides clinically meaningful prediction and characterization of CD19 CAR-T therapy response in B-cell lymphoma. The supervised model's attention weights enable identification of predictive TCR sequences, while unsupervised clustering reveals repertoire-level patterns associated with therapeutic efficacy. This dual approach combining high predictive performance with biological interpretability suggests potential for clinical translation in CAR-T patient stratification.

Keywords: T-cell receptor, deep learning, CAR-T therapy, CD19, B-cell lymphoma, multiple instance learning, variational autoencoder, attention mechanism, supervised learning, unsupervised learning

Contents

1	Introduction	4
1.1	CD19 CAR-T Cell Therapy and Clinical Challenge	4
1.2	CAR-T Product T-Cell Receptor Repertoires	4
1.3	T-Cell Receptor Biology and Repertoire Analysis	4
1.4	Deep Learning for TCR Repertoire Classification	5
1.5	Study Objectives	5
2	Methods	6
2.1	Study Design and Patient Cohort	6
2.2	TCR Sequencing and Data Processing	6
2.3	Feature Engineering and Encoding	7
2.4	DeepTCR Architecture	7
2.5	Training Configuration	8
2.6	Monte Carlo Cross-Validation	9
2.7	Statistical Analysis	9
2.8	Unsupervised Learning: VAE-Inspired Clustering	9
3	Results	10
3.1	Cohort Characteristics	10
3.2	Model Performance	11
3.3	Training Dynamics	12
3.4	Attention Analysis	12
3.5	Gene Usage Patterns	13
3.6	Top Predictive Sequences	14
3.7	Sequence Characteristics	15
3.8	Unsupervised Clustering Results	15
4	Discussion	16
4.1	Principal Findings	16
4.2	Comparison with Existing CAR-T Biomarkers	17
4.3	Biological Implications	17
4.4	Limitations	17
4.5	Future Directions	18
5	Conclusion	18

1 Introduction

1.1 CD19 CAR-T Cell Therapy and Clinical Challenge

Chimeric antigen receptor T-cell (CAR-T) therapy has revolutionized the treatment of hematological malignancies, particularly B-cell lymphomas [June et al., 2018]. CD19-targeted CAR-T products (axicabtagene ciloleucel, tisagenlecleucel, lisocabtagene maraleucel) have demonstrated remarkable efficacy in refractory/relapsed diffuse large B-cell lymphoma (DLBCL), follicular lymphoma, and other B-cell neoplasms [Neelapu et al., 2017, Schuster et al., 2019]. These therapies redirect autologous T cells to recognize and eliminate CD19-expressing malignant B cells, achieving complete response rates of 40–60% in heavily pretreated patients [Locke et al., 2019].

Despite these impressive outcomes, a substantial proportion of patients either fail to respond initially or relapse after initial response. Understanding the biological determinants of CAR-T therapy response remains a critical unmet need. Current predictive biomarkers include tumor burden, inflammatory markers (ferritin, CRP, LDH), and CAR-T product characteristics such as CD4/CD8 ratio and exhaustion markers [Locke et al., 2020, Deng et al., 2020]. However, these markers show variable predictive performance and fail to capture the complexity of T-cell repertoire composition within the infused product.

1.2 CAR-T Product T-Cell Receptor Repertoires

While CAR-T cells are engineered to express a synthetic chimeric receptor, they retain their endogenous T-cell receptor (TCR) repertoires [Fraietta et al., 2018]. These native TCRs reflect the clonal history, differentiation state, and functional potential of the underlying T cells. Single-cell studies have revealed that CAR-T products from responders show distinct TCR repertoire characteristics compared to non-responders, including differences in clonality, diversity, and specific TCR sequence features [Chen et al., 2021a]. Importantly, TCR repertoire analysis provides a window into T-cell fitness that complements CAR expression-based metrics.

1.3 T-Cell Receptor Biology and Repertoire Analysis

T-cell receptors serve as the primary molecular interface between the adaptive immune system and tumor cells. Each TCR is composed of an alpha and beta chain, with the complementarity-determining region 3 (CDR3) of each chain forming the primary contact surface with peptide-major histocompatibility complex (pMHC) molecules [Davis and Bjorkman, 1988, Rossjohn et al., 2015]. The CDR3 region arises through V(D)J recombination, creating a highly diverse repertoire estimated at 10^{15} to 10^{20} unique sequences in any individual [Robins et al., 2009].

The TCR beta chain (TRB) is particularly informative for repertoire analysis due to its higher diversity compared to the alpha chain and its dominant role in determining peptide specificity [Glanville et al., 2017]. TRB CDR3 sequences directly contact the peptide antigen, with specific amino acid motifs conferring recognition of distinct epitopes [Dash et al., 2017]. V-gene and J-gene usage patterns further constrain peptide recognition, with certain V-J combinations enriched in tumor-reactive T-cell populations [Zhang et al., 2020].

High-throughput sequencing has enabled comprehensive profiling of TCR repertoires from peripheral blood and tumor tissue [Robins et al., 2009, Mamedov et al., 2013]. Early studies established that increased TCR clonality and decreased diversity correlate with anti-PD-1 response in melanoma [Tumeh et al., 2014]. Subsequent work demonstrated that clonal expansion of tumor-infiltrating T cells predicts ipilimumab response [Postow et al., 2015].

1.4 Deep Learning for TCR Repertoire Classification

Traditional TCR repertoire analysis relies on aggregate statistics such as Shannon entropy, Simpson diversity, and clonality indices [Greiff et al., 2015]. While informative, these metrics collapse high-dimensional sequence data into scalar values, potentially discarding predictive information. Machine learning approaches, particularly deep learning, offer the capability to learn complex, non-linear patterns directly from raw sequence data [LeCun et al., 2015].

Multiple Instance Learning (MIL) provides an appropriate framework for repertoire classification [Dietterich et al., 1997]. In standard supervised learning, each sample has an associated label. However, TCR repertoires present a different challenge: each patient (label) is represented by thousands of sequences (instances), most of which are irrelevant to the classification task. MIL addresses this by treating the entire repertoire as a “bag” of sequences and learning to identify which sequences within each bag contribute to the patient-level prediction [Ilse et al., 2018].

Attention mechanisms, originally developed for natural language processing [Vaswani et al., 2017], provide biological interpretability to MIL models. The attention layer learns to assign importance weights to individual TCR sequences, effectively identifying tumor-reactive clonotypes without requiring sequence-level labels [Bahdanau et al., 2015].

DeepTCR, developed by Sidhom et al. [Sidhom et al., 2021], implements attention-based MIL specifically for TCR repertoire analysis. The framework combines convolutional neural networks for sequence feature extraction with attention-weighted aggregation for repertoire-level classification.

1.5 Study Objectives

The primary objective of this study was to develop and rigorously validate deep learning models for predicting CD19 CAR-T therapy response from pre-treatment TCR repertoires in refractory/relapsed B-cell lymphoma. We employed both **supervised** and **unsupervised** learning approaches to comprehensively characterize TCR repertoire features associated with therapeutic response.

Supervised Learning Objective: We hypothesized that attention-based multiple instance learning could identify predictive TCR signatures while maintaining biological interpretability through learned attention weights.

Unsupervised Learning Objective: We aimed to discover natural clustering patterns in TCR repertoires that might reveal underlying biological structure associated with CAR-T product fitness, independent of response labels.

Secondary objectives included: (1) quantifying model robustness through extensive Monte Carlo cross-validation, (2) comparing cluster distributions between responders and non-responders,

(3) characterizing TCR repertoire features associated with CAR-T response, and (4) evaluating the computational efficiency of deep learning approaches for clinical translation.

2 Methods

2.1 Study Design and Patient Cohort

This retrospective cohort study analyzed TCR repertoire data from 34 patients with refractory/relapsed B-cell lymphoma treated with CD19-targeted CAR-T cell therapy. Patients were classified as responders (n=18) or non-responders (n=16) based on clinical response assessment at day 90 post-infusion, with responders defined as achieving complete response (CR) or partial response (PR), and non-responders defined as stable disease (SD) or progressive disease (PD).

TCR sequencing was performed on the CAR-T cell products prior to infusion. The near-balanced cohort design (18 responders, 16 non-responders) helps prevent severe class imbalance from confounding model training. Patient and repertoire characteristics are summarized in Table 1.

Table 1: Patient and Repertoire Characteristics

Characteristic	All (n=34)	Responders (n=18)	Non-Responders (n=16)
<i>TCR Repertoire Metrics</i>			
Total sequences	239,637	130,924	108,713
Sequences/patient, mean	7,048	7,701	6,395
Sequences/patient, range	1,786–12,272	1,786–12,272	2,114–12,272
CDR3 length (AA), mean $\pm$ SD	14.45 $\pm$ 0.29	14.44 $\pm$ 0.25	14.46 $\pm$ 0.33
Unique V genes/patient	47.2	48.1	46.3
Unique J genes/patient	12.8	12.9	12.7

2.2 TCR Sequencing and Data Processing

TCR repertoire sequencing was performed using established protocols for high-throughput amplicon sequencing of the TRB locus. Total RNA was extracted from peripheral blood mononuclear cells (PBMCs), followed by reverse transcription and multiplex PCR amplification using primers targeting all functional TRBV and TRBJ gene segments. Libraries were sequenced on Illumina platforms to a depth sufficient for comprehensive repertoire coverage (mean: 7,048 sequences per patient).

The data processing pipeline is illustrated in Figure 1. Raw sequencing data underwent quality control including:

1. **Sequence alignment and annotation:** Reads were aligned to the human TRB reference genome (IMGT nomenclature) to identify V genes, J genes, and CDR3 boundaries.
2. **Quality filtering:** Sequences were retained if they contained in-frame CDR3 regions, no stop codons, complete V and J gene annotations, and CDR3 length ≤ 40 amino acids.
3. **Data standardization:** Column names were standardized to DeepTCR format requirements.

4. **Metadata integration:** Each sequence was annotated with response status at the patient level.




Figure 1: **Data Processing Pipeline.** Overview of the complete analytical workflow from raw TCR sequencing data through quality control, feature encoding, and DeepTCR model training. The pipeline processed 239,637 productive TRB sequences from 34 patients for analysis.

2.3 Feature Engineering and Encoding

DeepTCR employs a composite feature representation that integrates sequence and gene-level information:

Sequence Features: CDR3 amino acid sequences are encoded using a one-hot representation with a fixed maximum length of 40 amino acids. Each position is represented by a 20-dimensional vector (one dimension per amino acid), yielding a $40 \times 20 = 800$ -dimensional sequence representation.

Gene Features: V-gene and J-gene identities are encoded as separate one-hot vectors. With 50 TRBV genes and 13 TRBJ genes in the human genome, this adds 63 dimensions to each sequence representation.

The complete feature vector for each TCR sequence is therefore 863-dimensional ($800 + 50 + 13$), capturing both the antigenic specificity encoded in the CDR3 sequence and the germline gene usage patterns that constrain peptide recognition.

2.4 DeepTCR Architecture

DeepTCR implements an attention-based multiple instance learning architecture specifically designed for TCR repertoire classification (Figure 2). The model comprises four main components:

1. **Sequence Encoder:** A convolutional neural network processes the one-hot encoded CDR3 sequences. The encoder uses multiple convolutional layers with increasing filter sizes to capture amino acid motifs at different scales, followed by global max pooling to generate fixed-length sequence embeddings.

2. **Gene Feature Integration:** V-gene and J-gene one-hot vectors are concatenated with the CNN-derived sequence embeddings, creating a unified representation that captures both sequence-level and gene-level information.

3. Attention Mechanism: The attention layer computes importance weights for each TCR sequence within a patient’s repertoire using a two-layer neural network with tanh activation:

$$a_i = \frac{\exp(w^T \tanh(Vh_i))}{\sum_j \exp(w^T \tanh(Vh_j))} \quad (1)$$

where h_i is the embedding for sequence i , V and w are learned parameters, and the softmax normalization ensures attention weights sum to 1 per patient.

4. Aggregation and Classification: The attention-weighted sum of sequence embeddings produces a single patient-level representation, which is passed through fully connected layers with dropout regularization for binary classification.



Figure 2: DeepTCR Attention-Based Multiple Instance Learning Architecture

Figure 2: **DeepTCR Model Architecture.** The attention-based multiple instance learning framework processes individual TCR sequences through a convolutional encoder, integrates gene usage features, applies learned attention weights to identify important sequences, and aggregates to patient-level predictions.

2.5 Training Configuration

Model training was performed on NVIDIA H100 80GB GPU infrastructure using the following optimized hyperparameters:

- **Batch size:** 1024 sequences (optimized for H100 memory bandwidth)
- **Network size:** Large (increased hidden dimensions)
- **Learning rate:** Adam optimizer with default parameters
- **Epochs:** Minimum 10, with early stopping based on validation loss
- **Regularization:** Dropout with rate 0.5 in classification layers

2.6 Monte Carlo Cross-Validation

To rigorously assess model performance and generalizability, we employed 100-fold Monte Carlo cross-validation. In each fold:

1. Patients were randomly partitioned into training (75%) and test (25%) sets, maintaining class balance
2. The model was trained on the training set with early stopping
3. Performance was evaluated on the held-out test set using AUC-ROC

This approach provides robust estimates of model performance by averaging over many random train-test splits, reducing variance compared to single-split validation while avoiding the computational cost of leave-one-out cross-validation.

2.7 Statistical Analysis

Model performance was summarized using:

- **Mean AUC** with standard deviation across 100 folds
- **95% Confidence Interval:** Calculated as $\text{mean} \pm 1.96 \times \text{SE}$, where $\text{SE} = \text{SD}/\sqrt{n}$
- **Range and median** to characterize the distribution of fold-level performance

All analyses were performed using Python 3.10 with DeepTCR 2.0, TensorFlow 2.x, scikit-learn, and standard data science libraries.

2.8 Unsupervised Learning: VAE-Inspired Clustering

In addition to supervised prediction, we performed unsupervised analysis to discover natural patterns in TCR repertoires without using response labels. This approach is inspired by Variational Autoencoder (VAE) methodologies that have proven effective for biological sequence analysis [Kingma and Welling, 2014, Lopez et al., 2018].

Feature Extraction: TCR sequences were represented using k-mer frequency features ($k=3$), capturing local amino acid patterns within CDR3 regions. This representation yielded 6,213 unique 3-mer features across the dataset. V-gene and J-gene identities were encoded as one-hot vectors (49 V-genes, 13 J-genes), producing a combined feature matrix of 6,275 dimensions per sequence.

Dimensionality Reduction: Principal Component Analysis (PCA) was applied to reduce the high-dimensional feature space to a 50-dimensional latent representation, analogous to the latent space in VAE models. This latent space captures the primary modes of variation in TCR sequence composition and gene usage patterns.

Clustering: K-means clustering was applied to the latent space representations to identify natural groupings of TCR sequences. The number of clusters ($k=5$) was selected based on interpretability and cluster size distribution. Cluster assignments were analyzed for enrichment of responder versus non-responder sequences.

Visualization: t-distributed Stochastic Neighbor Embedding (t-SNE) was applied to visualize the high-dimensional latent space in two dimensions, enabling assessment of cluster separation and response status distribution.

3 Results

3.1 Cohort Characteristics

The study cohort comprised 34 patients with refractory/relapsed B-cell lymphoma treated with CD19 CAR-T cell therapy, with 18 responders and 16 non-responders. TCR repertoire sequencing of CAR-T products yielded 239,637 productive TRB sequences across all patients (Figure 3).

Figure 3: Cohort Characteristics and TCR Repertoire Overview




Figure 3: Cohort Overview. (A) Distribution of TCR sequences per patient by response status. (B) CDR3 length distribution across the cohort. (C) Summary of key cohort characteristics.

Repertoire sizes varied substantially across patients (range: 1,786–12,272 sequences per

patient), with responders showing slightly higher mean sequence counts (7,701 vs. 6,395). CDR3 length distributions were comparable between groups (mean: 14.45 ± 0.29 amino acids), consistent with productive, functional TCR sequences. V-gene and J-gene diversity was similar between responders and non-responders, with patients utilizing a mean of 47.2 unique V genes and 12.8 unique J genes. Per-patient repertoire characteristics and prediction confidence are detailed in Figure 9.

3.2 Model Performance

The DeepTCR model achieved robust predictive performance across 100 Monte Carlo cross-validation folds (Figure 4):

- **Mean AUC:** 0.754 ± 0.035 (SD)
- **Median AUC:** 0.765
- **Range:** 0.658–0.794
- **95% CI:** 0.747–0.761 (1,000 bootstrap iterations)

Figure 4: Model Performance - 100-Fold Monte Carlo Cross-Validation




Figure 4: Model Performance. (A) Receiver operating characteristic (ROC) curve showing mean performance (solid line) with 95% confidence band (shaded region) across 100 cross-validation folds. (B) Distribution of AUC values across folds, demonstrating consistent performance. (C) Summary statistics table.

The performance metrics indicate robust generalization across random train-test splits. Bootstrap analysis (1,000 iterations) confirmed a narrow 95% confidence interval of 0.747–0.761, providing statistical validation of model reliability. Complete performance metrics are provided in Table 2, with detailed AUC distribution analysis in Figure 8.

Table 2: Model Performance Summary

Metric	Value
Mean AUC	0.754
Standard Deviation	0.035
Median AUC	0.765
Minimum AUC	0.658
Maximum AUC	0.794
Range	0.136
95% Confidence Interval	0.747–0.761
Number of Folds	100
Bootstrap Iterations	1,000

3.3 Training Dynamics

Analysis of training dynamics across folds revealed consistent convergence behavior (Figure 5). Models typically achieved stable performance within 10 epochs, with early stopping preventing overfitting. The total training time for all 100 folds was 35 minutes on NVIDIA H100 GPU, demonstrating computational efficiency suitable for clinical deployment. Detailed computational performance metrics including GPU memory utilization are provided in Figure 11.

Figure 5: Training Dynamics Across 100 Folds




Figure 5: **Training Dynamics.** (A) Loss curves across training epochs showing convergence. (B) AUC progression during training. (C) Computational performance metrics.

3.4 Attention Analysis

The attention mechanism identified specific TCR sequences associated with immunotherapy response prediction (Figure 6). Analysis of 239,634 sequences revealed:

- **High-attention sequences:** 2,379 sequences (top 1%) with attention weights >99 th percentile
- **Attention distribution:** Mean attention weight of 1.42×10^{-4} , with maximum weight of 4.12×10^{-3}

- **Responder enrichment:** 72% of the top 100 highest-attention sequences originated from responder patients
- **Statistical significance:** Mann-Whitney U test confirmed significant difference in attention weights between responders and non-responders ($p < 10^{-178}$)




Figure 6: **Attention Analysis.** (A) Distribution of attention weights by response group showing responders (green) have distinct attention patterns compared to non-responders (red). (B) Comparison of mean attention weights per patient. (C) Statistical comparison with Mann-Whitney U test results. Complete attention weight distribution is provided in Figure 12.

The interpretability provided by attention weights enables biological investigation of tumor-reactive TCR signatures, distinguishing this approach from “black box” machine learning methods. The top-ranked sequence (CASRWEEDTEAFF, TRBV7-9) exhibited 29-fold higher attention than the population mean, suggesting strong predictive relevance for this specific clonotype. Detailed attention heatmaps across patients are provided in Figure 13.

3.5 Gene Usage Patterns

Analysis of V-gene and J-gene usage revealed differential patterns between responders and non-responders (Figure 7). Several V-gene families showed significant enrichment in responder repertoires, potentially reflecting shared tumor antigen recognition. Detailed J-gene usage analysis is provided in Figure 10.




Figure 7: **Gene Usage Patterns.** (A) V-gene usage frequency by response status. (B) J-gene usage frequency by response status. (C) V-J pairing preferences in responders vs. non-responders.

3.6 Top Predictive Sequences

Enrichment analysis of high-attention sequences identified specific V-gene and J-gene usage patterns associated with immunotherapy response (Figure 14). Key findings include:

- **V-gene enrichment:** TRBV5-8 showed the highest enrichment in high-attention sequences (1.54-fold, $p = 0.23$), followed by TRBV30 (1.45-fold) and TRBV5-4 (1.44-fold)
- **J-gene enrichment:** TRBJ2-2 exhibited 1.17-fold enrichment in high-attention sequences, followed by TRBJ2-7 (1.13-fold)
- **V-J pairing:** The most frequent V-J combination in high-attention sequences was TRBV28+TRBJ2-7 (42 sequences), followed by TRBV20-1+TRBJ2-7 (40 sequences)

The top 10 highest-attention sequences are listed in Table 3. Notably, 8 of the top 10 sequences originated from responder patients, with the highest-attention sequence (CASR-WEEDTEAFF) utilizing TRBV7-9 and TRBJ1-1.

Table 3: **Top 10 Predictive TCR Sequences by Attention Weight**

Rank	CDR3 Sequence	V Gene	J Gene	Response	Attention
1	CASRWEEDTEAFF	TRBV7-9	TRBJ1-1	R	4.12×10^{-3}
2	CASSRGGDFYNEQFF	TRBV3-1	TRBJ2-1	R	3.97×10^{-3}
3	CASSLGRTYEQYF	TRBV27	TRBJ2-7	R	3.69×10^{-3}
4	CASSVTSYNEQFF	TRBV9	TRBJ2-1	NR	3.61×10^{-3}
5	CASSLGGETQYF	TRBV11-2	TRBJ2-5	R	3.40×10^{-3}
6	CASSRGGDFYNEQFF	TRBV3-1	TRBJ2-1	R	3.30×10^{-3}
7	CASSLRFGGEQYF	TRBV27	TRBJ2-7	R	3.29×10^{-3}
8	CASSFRITYNEQFF	TRBV27	TRBJ2-1	R	3.06×10^{-3}
9	CASTVGERQYF	TRBV6-2	TRBJ2-7	R	3.05×10^{-3}
10	CASTPNGAVAEAFF	TRBV25-1	TRBJ1-1	NR	2.99×10^{-3}

3.7 Sequence Characteristics

Analysis of CDR3 sequence characteristics revealed subtle but statistically significant differences between high-attention and low-attention sequences (Figure 16):

- **CDR3 length:** High-attention sequences showed slightly longer CDR3 regions (14.52 ± 1.80 AA) compared to low-attention sequences (14.46 ± 1.78 AA), though this difference was modest
- **Amino acid composition:** Tryptophan (W) and valine (V) showed the highest enrichment in high-attention sequences (1.15-fold and 1.13-fold, respectively)
- **Responder vs. non-responder:** CDR3 length distributions differed significantly between groups (Kolmogorov-Smirnov test, $p < 10^{-31}$; responders: 14.49 ± 1.75 AA, non-responders: 14.42 ± 1.83 AA)

These findings suggest that while CDR3 length contributes minimally to prediction, specific amino acid motifs and sequence patterns captured by the deep learning model carry predictive information for CAR-T therapy response.

3.8 Unsupervised Clustering Results

The unsupervised VAE-inspired analysis revealed distinct clustering patterns in the TCR repertoire that partially segregate responders from non-responders (Figure 17). Key findings include:

- **Cluster identification:** Five distinct TCR clusters were identified based on k-mer sequence features and gene usage patterns
- **Responder-enriched cluster:** Cluster 4 showed the highest responder enrichment (1.82-fold), containing 99.3% responder-derived sequences
- **Cluster 2 enrichment:** This cluster showed moderate responder enrichment (1.05-fold, 57.5% responder sequences)

- **Dominant cluster:** Cluster 3 contained the majority of sequences (52.9%) with near-equal distribution between response groups

Patient-level analysis revealed that responders showed higher representation in Cluster 2 (21.9% vs. 19.4% in non-responders) and exclusive representation in Cluster 4 (0.3% vs. 0.0%). These differences, while modest at the overall level, suggest that specific TCR subpopulations may be associated with CAR-T product fitness.

Gene usage patterns by cluster revealed distinct V-J preferences:

- Cluster 0: Enriched for TRBV20-1, TRBV7-9, TRBV9 with TRBJ2-1, TRBJ2-7
- Cluster 2: Enriched for TRBV12-4, TRBV4-1 with TRBJ1-1, TRBJ2-5
- Cluster 4 (responder-enriched): Dominated by TRBV3-1 with TRBJ1-3

The convergence of supervised and unsupervised findings—both identifying TRBV3-1 as a distinguishing feature—provides independent validation of this gene’s association with CAR-T response.

4 Discussion

4.1 Principal Findings

This study demonstrates that deep learning analysis of TCR repertoires from CD19 CAR-T cell products can predict therapeutic response in refractory/relapsed B-cell lymphoma with clinically meaningful accuracy. Using both **supervised** and **unsupervised** approaches, we identified TCR signatures associated with CAR-T efficacy.

Supervised Learning: The attention-based MIL model achieved an AUC of 0.754 ± 0.035 (SD) across 100 Monte Carlo folds, with a bootstrap-derived 95% confidence interval of 0.747–0.761. This performance compares favorably with existing CAR-T response biomarkers such as CD4/CD8 ratio, exhaustion markers, and inflammatory indices [Deng et al., 2020, Locke et al., 2020].

Unsupervised Learning: VAE-inspired clustering revealed five distinct TCR populations, with Cluster 4 showing 1.82-fold responder enrichment (99.3% responder-derived sequences). The convergence of supervised attention analysis and unsupervised clustering—both identifying TRBV3-1 as a distinguishing feature—provides independent validation of gene usage associations with CAR-T response.

Several aspects of our findings merit emphasis. First, the supervised model achieved consistent performance across 100 independent cross-validation folds, suggesting that DeepTCR learned generalizable patterns rather than overfitting. Second, the attention mechanism provides biological interpretability, identifying 2,379 high-attention sequences with 72% of the top 100 originating from responder patients. Third, the unsupervised analysis revealed natural repertoire structure that partially segregates response groups without using labels, suggesting that TCR composition reflects underlying T-cell fitness properties relevant to CAR-T efficacy.

4.2 Comparison with Existing CAR-T Biomarkers

Current predictive biomarkers for CAR-T response include tumor burden, inflammatory markers (LDH, ferritin, CRP), and product characteristics such as CD4/CD8 ratio and exhaustion phenotypes [Locke et al., 2020, Deng et al., 2020]. Fraietta et al. demonstrated that CAR-T products with memory phenotypes and lower exhaustion markers show improved persistence and efficacy [Fraietta et al., 2018].

Our TCR-based approach offers complementary information. While existing biomarkers focus on surface phenotypes and metabolic states, TCR repertoire analysis captures the clonal composition and antigen-recognition potential of the T-cell pool. The combination of TCR features with existing biomarkers may improve predictive accuracy.

4.3 Biological Implications

The success of both supervised and unsupervised approaches supports the hypothesis that predictive information is encoded in TCR repertoire structure. For CAR-T therapy, this has several implications:

T-cell fitness: High-attention TCR sequences and responder-enriched clusters may mark T-cell populations with favorable properties for CAR-T manufacturing and persistence. The endogenous TCR repertoire reflects the clonal history and differentiation state of T cells, which influences CAR-T product quality [Fraietta et al., 2018].

Convergent selection: The enrichment of specific V-genes (TRBV3-1, TRBV5-8) in both high-attention sequences and responder-enriched clusters suggests convergent selection for T cells with particular antigenic specificities. These may represent T-cell populations with favorable properties for adoptive cell therapy.

Product heterogeneity: The identification of distinct TCR clusters within CAR-T products highlights the heterogeneity of these cellular therapies. Understanding which TCR subpopulations contribute to efficacy could inform manufacturing optimization.

4.4 Limitations

Several limitations of this study warrant consideration. First, the sample size of 34 patients, while sufficient for proof-of-principle, limits statistical power and generalizability. Validation in larger, independent CAR-T cohorts is essential before clinical implementation.

Second, the cohort design (18 responders, 16 non-responders) may not reflect real-world CAR-T response rates, which vary by disease subtype and patient characteristics. Model calibration should be evaluated in populations with varying class distributions.

Third, our unsupervised clustering approach, while inspired by VAE methodology, uses PCA-based dimensionality reduction rather than true variational inference. Future work should implement full VAE architectures for potentially improved latent representations.

Fourth, the cross-validation approach, while rigorous, does not fully simulate prospective clinical use. Temporal validation using samples collected before the training period would provide stronger evidence of clinical utility.

4.5 Future Directions

Several extensions of this work are warranted:

Multi-center validation: Studies with larger sample sizes would establish generalizability across patient populations, CAR-T products, and sequencing platforms.

Biomarker integration: Combining TCR features with existing CAR-T biomarkers (CD4/CD8 ratio, exhaustion markers, tumor burden) may improve predictive accuracy.

Longitudinal analysis: Tracking TCR repertoire evolution during CAR-T expansion and persistence could enable dynamic response prediction.

Functional validation: High-attention TCR sequences identified by supervised learning and responder-enriched clusters from unsupervised analysis should be functionally characterized to understand their role in CAR-T efficacy.

Manufacturing optimization: If specific TCR subpopulations correlate with response, this could inform CAR-T manufacturing protocols to enrich for favorable T-cell populations.

5 Conclusion

We have demonstrated that deep learning analysis of TCR repertoires from CD19 CAR-T cell products provides clinically meaningful prediction of therapeutic response in refractory/relapsed B-cell lymphoma. Using complementary **supervised** and **unsupervised** approaches, we identified TCR signatures associated with CAR-T efficacy.

Supervised learning achieved an AUC of 0.754 ± 0.035 (SD) across 100 Monte Carlo cross-validation folds (95% CI: 0.747–0.761), with attention weights identifying 2,379 high-attention sequences. Notably, 72% of the top 100 predictive sequences originated from responder patients.

Unsupervised learning revealed five distinct TCR clusters, with Cluster 4 showing 1.82-fold responder enrichment (99.3% responder-derived). The convergence of both approaches on TRBV3-1 as a distinguishing feature provides independent validation of gene usage associations.

This dual approach combines four desirable properties for clinical translation: (1) predictive accuracy comparable to existing CAR-T biomarkers, (2) biological interpretability through attention-based sequence identification, (3) unsupervised discovery of repertoire structure without requiring labels, and (4) computational efficiency enabling real-time prediction. These findings support further development of TCR-based biomarkers for CAR-T patient stratification, with validation in larger cohorts as the critical next step toward clinical implementation.

Acknowledgments

M.J. extends heartfelt gratitude to **Dr. Sherry Bhalla** for her invaluable guidance, mentorship, and continuous support throughout this research project. Her expertise in computational biology and dedication to training the next generation of researchers made this work possible. We also thank all members of the CSIR-IGIB laboratory for their assistance, insightful discussions, and collaborative spirit during this training period.

We gratefully acknowledge the **Lambda Labs GPU Grant Program** for providing access to NVIDIA H100 80GB GPU infrastructure, which was essential for training the deep learning

models in this study.

This work was supported by the **Council of Scientific and Industrial Research (CSIR)** and the **Institute of Genomics and Integrative Biology (IGIB)**, New Delhi, India. M.J. acknowledges support as a Research Trainee in the Deep Learning/Machine Learning for Genomic Analysis program at CSIR-IGIB.

We acknowledge the use of AI-assisted development tools during this research: **Claude Code** [Anthropic, 2025], an agentic coding tool by Anthropic, and **OpenAI Codex** [Chen et al., 2021b] were used to assist with code development, debugging, and optimization of the analysis pipeline.

We thank the developers of the **DeepTCR** framework [Sidhom et al., 2021] for making their software openly available. M.J. contributed to the DeepTCR project by assisting with the upgrade to Python 3.10 and TensorFlow 2.12 (version 2.1.28), as acknowledged in the official DeepTCR repository (<https://github.com/sidhomj/DeepTCR>).

We thank the patients who contributed samples to this study, without whom this research would not have been possible.

Conflict of Interest

The authors declare no competing interests.

Code and Data Availability

All code, trained models, and analysis scripts developed in this study are publicly available on GitHub:

https://github.com/MateenahJAHAN/Final_Deliverable_DeepTCRResearch_Mateenah_Jahan_CSIR

The repository includes:

- Complete analysis pipeline (Python scripts)
- Trained model checkpoints from 100-fold Monte Carlo cross-validation
- Figure generation scripts
- LaTeX source files for this manuscript

The DeepTCR framework is available at <https://github.com/sidhomj/DeepTCR>. The processed TCR dataset used in this study is included in the repository.

References

Anthropic. Claude code: Agentic coding tool for software development. <https://docs.anthropic.com/en/docs/claude-code>, 2025. Accessed: December 2025.

- Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural machine translation by jointly learning to align and translate. *arXiv preprint arXiv:1409.0473*, 2015. URL <https://arxiv.org/abs/1409.0473>. Published at ICLR 2015.
- Grace M Chen, Changya Chen, Rajan K Das, Peng Gao, Chun-Hui Chen, Sze Shun Kuo, Danielle L Gibbons, Brett E Bernstein, Dhrubajyoti Bhattacharya, Vikram Bhoje, et al. Single-cell profiling of t-cell repertoire and phenotype during CAR-t therapy for b-cell malignancies. *Blood*, 137(8):1072–1083, 2021a. doi: 10.1182/blood.2020006568.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Pondé de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021b. doi: 10.48550/arXiv.2107.03374. URL <https://arxiv.org/abs/2107.03374>.
- Pradyot Dash, Andrew J Fiore-Gartland, Tomer Hertz, George C Wang, Shalini Sharma, Aisha Souquette, Jeremy Chase Crawford, E Bridie Clemens, Thi HO Nguyen, Katherine Kedzierska, et al. Quantifiable predictive features define epitope-specific t cell receptor repertoires. *Nature*, 547(7661):89–93, 2017. doi: 10.1038/nature22383.
- Mark M Davis and Pamela J Bjorkman. T-cell antigen receptor genes and t-cell recognition. *Nature*, 334(6181):395–402, 1988. doi: 10.1038/334395a0.
- Qing Deng, Gang Han, Nadia Puebla-Osorio, Michelle CJ Ma, Paolo Strati, Bagdon Chaber, Veronica Saez-Ibarra, Diana Santos, Ashley Orozco, Sapna Patel, et al. Characteristics of anti-CD19 CAR t cell infusion products associated with efficacy and toxicity in patients with large b cell lymphomas. *Nature Medicine*, 26(12):1878–1887, 2020. doi: 10.1038/s41591-020-1061-7.
- Thomas G Dietterich, Richard H Lathrop, and Tomás Lozano-Pérez. Solving the multiple instance problem with axis-parallel rectangles. *Artificial Intelligence*, 89(1-2):31–71, 1997. doi: 10.1016/S0004-3702(96)00034-3.
- Joseph A Fraietta, Simon F Lacey, Elena J Orlando, Iulian Pruteanu-Malinici, Mercy Gohil, Stefan Lundh, Alina C Boesteanu, Yong Wang, Roddy S O’Connor, Wei-Ting Hwang, et al. Determinants of response and resistance to CD19 chimeric antigen receptor (CAR) t cell therapy of chronic lymphocytic leukemia. *Nature Medicine*, 24(5):563–571, 2018. doi: 10.1038/s41591-018-0010-1.
- Jacob Glanville, Huang Huang, Allison Nau, Olivia Hatton, Lisa E Wagar, Florian Rubelt, Xuhuai Ji, Arnold Han, Sheri M Krams, Christina Pettus, et al. Identifying specificity groups in the t cell receptor repertoire. *Nature*, 547(7661):94–98, 2017. doi: 10.1038/nature22976.
- Victor Greiff, Enkelejda Miho, Ulrike Menzel, and Sai T Reddy. Bioinformatic and statistical analysis of adaptive immune repertoires. *Trends in Immunology*, 36(11):738–749, 2015. doi: 10.1016/j.it.2015.09.006.
- Maximilian Ilse, Jakub Tomczak, and Max Welling. Attention-based deep multiple instance learning. In *International Conference on Machine Learning*, volume 80 of *Proceedings of*

- Machine Learning Research*, pages 2127–2136. PMLR, 2018. URL <https://proceedings.mlr.press/v80/ilse18a.html>.
- Carl H June, Roddy S O’Connor, Omkar U Kawalekar, Saba Ghassemi, and Michael C Milone. CAR t cell immunotherapy for human cancer. *Science*, 359(6382):1361–1365, 2018. doi: 10.1126/science.aar6711.
- Diederik P Kingma and Max Welling. Auto-encoding variational Bayes. *arXiv preprint arXiv:1312.6114*, 2014. URL <https://arxiv.org/abs/1312.6114>. Published at ICLR 2014.
- Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *Nature*, 521(7553):436–444, 2015. doi: 10.1038/nature14539.
- Frederick L Locke, Armin Ghobadi, Caron A Jacobson, David B Miklos, Lazaros J Lekakis, Olalekan O Oluwole, Yi Lin, Ira Braunschweig, Brian T Hill, John M Timmerman, et al. Long-term safety and activity of axicabtagene ciloleucel in refractory large b-cell lymphoma (ZUMA-1): a single-arm, multicentre, phase 1–2 trial. *The Lancet Oncology*, 20(1):31–42, 2019. doi: 10.1016/S1470-2045(18)30864-7.
- Frederick L Locke, John M Rossi, Sattva S Neelapu, Caron A Jacobson, David B Miklos, Armin Ghobadi, Olalekan O Oluwole, Patrick M Reagan, Lazaros J Lekakis, Yi Lin, et al. Tumor burden, inflammation, and product attributes determine outcomes of axicabtagene ciloleucel in large b-cell lymphoma. *Blood Advances*, 4(19):4898–4911, 2020. doi: 10.1182/bloodadvances.2020002394.
- Romain Lopez, Jeffrey Regier, Michael B Cole, Michael I Jordan, and Nir Yosef. Deep generative modeling for single-cell transcriptomics. *Nature Methods*, 15(12):1053–1058, 2018. doi: 10.1038/s41592-018-0229-2.
- Ilgar Z Mamedov, Olga V Britanova, Ivan V Zvyagin, Maria A Turchaninova, Dmitriy A Bolotin, Ekaterina V Putintseva, Yuri B Lebedev, and Dmitriy M Chudakov. Preparing unbiased t-cell receptor and antibody cDNA libraries for the deep next generation sequencing profiling. *Frontiers in Immunology*, 4:456, 2013. doi: 10.3389/fimmu.2013.00456.
- Sattva S Neelapu, Frederick L Locke, Nancy L Bartlett, Lazaros J Lekakis, David B Miklos, Caron A Jacobson, Ira Braunschweig, Olalekan O Oluwole, Tanya Siddiqi, Yi Lin, et al. Axicabtagene ciloleucel CAR t-cell therapy in refractory large b-cell lymphoma. *New England Journal of Medicine*, 377(26):2531–2544, 2017. doi: 10.1056/NEJMoa1707447.
- Michael A Postow, Mariana Manuel, Phillip Wong, Jianda Yuan, Zhiqiang Dong, Cailian Liu, Stephanie Perez, Isabelle Tanneau, Magalie Noel, Anne Courtier, et al. Peripheral t cell receptor diversity is associated with clinical outcomes following ipilimumab treatment in metastatic melanoma. *Journal for ImmunoTherapy of Cancer*, 3(1):23, 2015. doi: 10.1186/s40425-015-0070-4.
- Harlan S Robins, Paulo V Campregher, Santosh K Srivastava, Abigail Wachter, Cameron J Turtle, Orsalem Kahsai, Stanley R Riddell, Edus H Warren, and Christopher S Carlson.

Comprehensive assessment of t-cell receptor β -chain diversity in $\alpha\beta$ t cells. *Blood*, 114(19): 4099–4107, 2009. doi: 10.1182/blood-2009-04-217604.

Jamie Rossjohn, Stephanie Gras, John J Miles, Stephen J Turner, Dale I Godfrey, and James McCluskey. T cell antigen receptor recognition of antigen-presenting molecules. *Annual Review of Immunology*, 33:169–200, 2015. doi: 10.1146/annurev-immunol-032414-112334.

Stephen J Schuster, Michael R Bishop, Constantine S Tam, Edmund K Waller, Peter Borchmann, Joseph P McGuirk, Ulrich Jäger, Samantha Jaglowski, Charalambos Andreadis, Jason R Westin, et al. Tisagenlecleucel in adult relapsed or refractory diffuse large b-cell lymphoma. *New England Journal of Medicine*, 380(1):45–56, 2019. doi: 10.1056/NEJMoa1804980.

John-William Sidhom, H Benjamin Larman, Drew M Pardoll, and Alexander S Baras. DeepTCR is a deep learning framework for revealing sequence concepts within t-cell repertoires. *Nature Communications*, 12(1):1605, 2021. doi: 10.1038/s41467-021-21879-w.

Paul C Tumeh, Christina L Harview, Jennifer H Yearley, I Peter Shintaku, Emma JM Taylor, Lidia Robert, Bartosz Chmielowski, Marko Spasic, Gina Henry, Voicu Ciobanu, et al. PD-1 blockade induces responses by inhibiting adaptive immune resistance. *Nature*, 515(7528): 568–571, 2014. doi: 10.1038/nature13954.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, pages 5998–6008, 2017. URL <https://papers.nips.cc/paper/7181-attention-is-all-you-need>.

Shu-Qi Zhang, Ke-Yue Ma, Alexandra A Schonnesen, Mingliang Zhang, Chenglong He, Eric Sun, Chad M Williams, Anish Jain, Biplab Bhattacharya, Victoria Fang-Landau, et al. Investigation of antigen-specific t-cell receptor clusters in human cancers. *Clinical Cancer Research*, 26(6):1359–1371, 2020. doi: 10.1158/1078-0432.CCR-19-3249.

Supplementary Information

Figure S1: Detailed AUC Analysis




Figure 8: **Detailed AUC Analysis.** (A) Histogram of AUC values across 100 folds. (B) Box plot with individual fold values. (C) Fold-by-fold AUC stability analysis.

Figure S2: Per-Patient Analysis




Figure 9: **Per-Patient Analysis.** (A) Repertoire size by patient. (B) Prediction confidence by patient. (C) Attention weight distributions across patients.

Figure S3: J-Gene Analysis




Figure 10: **J-Gene Usage Analysis.** (A) J-gene frequency comparison between responders and non-responders. (B) J-gene diversity metrics.

Figure S4: Computational Performance


Figure 11: **Computational Performance.** (A) Training time per fold. (B) GPU memory utilization. (C) Throughput analysis on H100 GPU.

Figure S5: Attention Weight Distribution


Figure 12: **Attention Weight Distribution.** Complete distribution of attention weights across all 239,634 TCR sequences. The log-scale histogram reveals a highly skewed distribution with most sequences receiving near-zero attention, while a small subset (top 1%, $n=2,379$) exhibit high attention weights indicative of predictive importance.

Figure S6: Patient-Level Attention Heatmap

Figure 13: Patient-Level Attention Heatmap. Heatmap visualization of attention weight statistics across 34 patients. Rows represent individual patients, with color intensity indicating mean attention weight. Patients are grouped by response status (responders vs. non-responders).

Figure S7: Top Predictive Sequences Analysis



Figure 14: **Top Predictive Sequences Analysis.** (A) V-gene enrichment in top 100 high-attention sequences compared to background frequency. (B) J-gene enrichment analysis. (C) Response distribution showing 72% of top 100 sequences originate from responder patients. (D) CDR3 length comparison between top sequences and overall population.


Figure S8: Responder vs. Non-Responder Comparison**Figure: Responder vs Non-Responder Comparison**

Figure 15: Responder vs. Non-Responder Statistical Comparison. (A) Attention weight distributions by response group with Mann-Whitney U test results ($p < 10^{-178}$). (B) CDR3 length distributions with Kolmogorov-Smirnov test ($p < 10^{-31}$). (C) V-gene usage comparison. (D) High-attention sequence counts by group.

Figure S9: Sequence Characteristics Analysis



Figure 16: **Sequence Characteristics Analysis.** (A) CDR3 length distribution comparing high-attention vs. low-attention sequences. (B) Amino acid composition enrichment in high-attention sequences. (C) Amino acid property distribution (hydrophobic, polar, charged). (D) Top V-J gene pairings in high-attention sequences.


Figure S10: Unsupervised Clustering Analysis

Figure 17: Unsupervised Clustering Analysis. (A) t-SNE visualization of TCR repertoire colored by cluster assignment ($k=5$ clusters identified). (B) t-SNE visualization colored by response status (green: responder, red: non-responder). The unsupervised analysis reveals natural groupings in TCR repertoire that partially segregate by response status. Cluster 4 (0.3% of sequences) shows 1.82-fold responder enrichment with 99.3% of sequences from responder patients, suggesting this TCR subpopulation may mark T cells with favorable properties for CAR-T efficacy.